

TOKEN COST REPORT

What one token actually costs, tool by tool.

Every figure below is measured, not estimated — pulled directly from `token_usage`, the row OpenAI's own `completion.usage` writes on every real generation. No projections, no guesswork.

BLENDDED AVG / TOKEN \$0.0000058/tok ≈ \$5.79 per million tokens, across all 42 tools	TOTAL TOKENS LOGGED 299,794 across 159 generations	TOTAL SPEND LOGGED \$1.7369 real OpenAI billing, all tools	COST / GENERATION AVG \$0.0109 median tool lands near 1¢
---	--	--	--

VERIFIED AGAINST THE OPENAI BILLING CONSOLE — ONE CORRECTION

Reconciling this data against a real OpenAI console export (Jul 4 – Aug 3) confirms every **text** tool below to the cent — logged cost and console cost match almost exactly. **AI-generated images do not**: Jooma's internal price list charges `$0.0155` per image, but the console shows OpenAI actually billing **~\$0.0455 per image** that day — a **~2.9× understatement**. This only affects tools that generate AI images (chiefly the **Slideshow Generator** in AI-image mode); every other figure on this page is unaffected. Corrected slideshow cost: see the callout below the table.

Per-tool breakdown

Sorted by cost per token, highest first. This is the number that matters for pricing packages — it's what stays constant regardless of how long a given generation happens to run. `outlier` flags tools running meaningfully above the pack; `efficient` flags the cheapest. **Token-based costs only** — see the image-cost correction above and below for tools that also generate images.

TOOL	\$ / TOKEN	\$ / 1M TOKENS	AVG TOKENS / GEN	AVG \$ / GEN	RELATIVE
Eyfs Planner outlier 2 runs	\$0.0000088	\$8.84	6,224	\$0.055035	
Targeted Intervention outlier 2 runs	\$0.0000076	\$7.65	1,551	\$0.011856	
Sensory Activities outlier 2 runs	\$0.0000076	\$7.63	2,778	\$0.021199	
School Improvement Plan 3 runs	\$0.0000073	\$7.28	2,357	\$0.017152	
Performance Management 2 runs	\$0.0000070	\$7.04	1,561	\$0.010982	
Topic Overview 2 runs	\$0.0000069	\$6.87	2,128	\$0.014620	
Comprehension Generator 5 runs	\$0.0000068	\$6.81	1,312	\$0.008937	
Meeting Planner 4 runs	\$0.0000068	\$6.81	1,745	\$0.011884	
Ect Report Writer 2 runs	\$0.0000066	\$6.63	2,192	\$0.014526	
Policy Generator 2 runs	\$0.0000066	\$6.61	1,459	\$0.009648	
Assembly Planner 2 runs	\$0.0000066	\$6.60	1,675	\$0.011062	
Lesson Observation Report 2 runs	\$0.0000066	\$6.56	1,606	\$0.010539	
Worksheet Generator 4 runs	\$0.0000065	\$6.50	1,858	\$0.012074	
Eyfs Action Plan 2 runs	\$0.0000065	\$6.49	2,762	\$0.017919	
Homework Generator 2 runs	\$0.0000065	\$6.47	1,633	\$0.010561	
Inspection Prep 2 runs	\$0.0000065	\$6.47	3,541	\$0.022892	

TOOL	\$ / TOKEN	\$ / 1M TOKENS	AVG TOKENS / GEN	AVG \$ / GEN	RELATIVE
Model Text Generator 2 runs	\$0.0000064	\$6.43	1,519	\$0.009771	
Exam Question Generator 3 runs	\$0.0000064	\$6.38	1,219	\$0.007773	
Medium Term Planner 3 runs	\$0.0000063	\$6.32	1,404	\$0.008881	
Lesson Planner 10 runs	\$0.0000063	\$6.31	3,304	\$0.020829	
Cover Lesson 3 runs	\$0.0000063	\$6.30	1,747	\$0.011003	
Newsletter Writer 2 runs	\$0.0000063	\$6.27	937	\$0.005874	
Risk Assessment 6 runs	\$0.0000062	\$6.17	1,575	\$0.009714	
Generate Lesson Outline 5 runs	\$0.0000061	\$6.10	437	\$0.002670	
Pupil Premium Planner 2 runs	\$0.0000061	\$6.09	1,222	\$0.007443	
Learning Walk Report 2 runs	\$0.0000059	\$5.93	1,504	\$0.008919	
Model Answer Generator 2 runs	\$0.0000059	\$5.89	2,062	\$0.012146	
Phonics Support 3 runs	\$0.0000059	\$5.86	3,304	\$0.019371	
Quiz Generator 2 runs	\$0.0000058	\$5.80	869	\$0.005033	
Behaviour Support Plan 3 runs	\$0.0000055	\$5.48	3,261	\$0.017877	
Generate Audio 3 runs	\$0.0000052	\$5.21	693	\$0.003608	
Smart Targets 3 runs	\$0.0000051	\$5.07	1,160	\$0.005878	
One Page Profile 3 runs	\$0.0000050	\$5.03	1,366	\$0.006875	

TOOL	\$ / TOKEN	\$ / 1M TOKENS	AVG TOKENS / GEN	AVG \$ / GEN	RELATIVE
Cpd Slideshow 2 runs	\$0.0000050	\$4.99	1,862	\$0.009285	
Letter Writer 3 runs	\$0.0000047	\$4.73	576	\$0.002726	
Generate Slideshow 37 runs	\$0.0000047	\$4.67	2,639	\$0.012323	
Report Writer 2 runs	\$0.0000041	\$4.11	918	\$0.003773	
Generate Activity 2 runs	\$0.0000041	\$4.08	455	\$0.001853	
Suggest Vocabulary 6 runs	\$0.0000038	\$3.84	173	\$0.000666	
Find Youtube 2 runs	\$0.0000032	\$3.16	440	\$0.001388	
Edit Text 2 runs efficient	\$0.0000002	\$0.24	254	\$0.000062	
Suggest Subject efficient 6 runs	\$0.0000002	\$0.17	316	\$0.000052	

Slideshow Generator — corrected image cost

The only tool where images meaningfully affect total cost. Text and audio components (top two rows) are token-based and already accurate; the AI-image component is corrected below using the real ~\$0.0455/image rate confirmed against the OpenAI console, in place of the \$0.0155 flat rate currently in the code.

MODE	TEXT + AUDIO	IMAGES (AS LOGGED)	IMAGES (CORRECTED)	TOTAL (AS REPORTED)	TOTAL (CORRECTED)
Web / no AI images	\$0.0401	\$0.00	\$0.00	\$0.0401	
AI images avg of 4 decks	~\$0.045	~\$0.13	~\$0.38	~\$0.18	

A fully AI-illustrated deck costs roughly **2.5× more than previously reported** — closer to 45¢ than 18¢ per deck. The web/Pixabay mode is untouched by this correction and remains the cheap default at ~4¢.

HOW THIS WAS MEASURED

Every text-generation route in Jooma calls OpenAI through a shared helper (`streamChat()`), which reads `completion.usage` off the real API response — the exact count OpenAI bills on — and writes one row to `token_usage` per generation. Nothing in the table above is a token-counter estimate.

Cost per token = total logged `cost_usd` ÷ total logged tokens (prompt + completion), per tool. The blended average weights every tool by its actual token volume, not a flat mean across tools. **Table figures cover text generation only** — AI image cost is priced separately (see the correction above) and isn't token-based.

READING THE NUMBERS

- All 42 tools cost within a **~2.5×** band per token (\$0.00000017–\$0.0000088) — gpt-4o pricing is flat, so this spread comes almost entirely from each tool's *prompt-to-output ratio*, not the model.
- `eyfs-planner` is the clear outlier — long system prompt relative to a comparatively short output pulls its per-token cost highest.
- `suggest-subject` and `edit-text` sit near-zero — both run on `gpt-4o-mini` or with minimal output.
- Sample sizes are thin for many tools (2–3 runs) — treat this as directionally solid, not statistically final. More production usage will sharpen it.